

Big Compute



El héroe para nuestros problemas
de computación masiva

Jose Angel Fernandez - @jangelandez

Iria Quiroga - @iriaq

Iria Quiroga

Service Delivery Manager @  **Microsoft**



@iriaq



iriaquiroya



Jose angel Fernandez

Cloud Solution Architect @  **Microsoft**



@jangelandez



jangelandez



Patrocinadores Locales



Colabora



Agenda

- HPC on Azure, seriously?
- The key elements
 - Storage
 - Networking
 - Compute
 - Management
- Let's play



HPC on Azure Seriously?

What people initially
thinks when we
discuss about
running High
Performance
Applications on the
Cloud in general?



Petascade computing for Terapixel Visualization

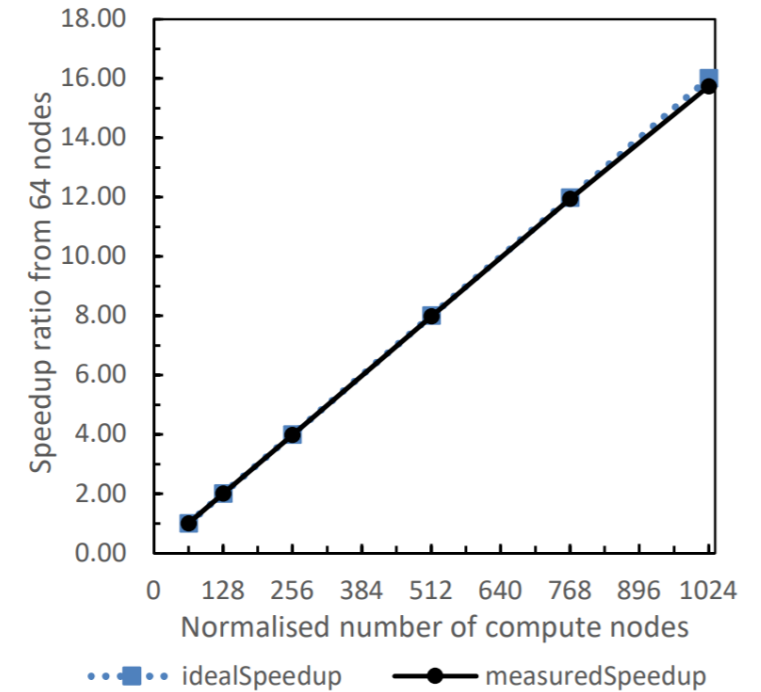


Visualization results:

- The first terapixel visualization of an urban digital twin, photo-realistic rendering using Blender Cycles path tracing.
- Tiled image output format makes supercomputer visualization outputs accessible on low cost, thin client devices.
- **Rendering time reduced from an estimated 34 days on one V100 GPU to 48 min. on 1024 V100 GPU.**

Visual supercomputing results:

- Efficiently deployed a 1024 GPU, 14 PFlop cloud supercomputer: faster than any public GPU HPC system in the UK.
- **Cost for one performance scaling graph: <£20,000, providing access to a > £10 million supercomputer.**
- **Azure upgrade** (from K80 to V100) during the project, transparently benefited from a 3x performance increase.
- Enables experiments with future scaled performance approximately 20 years ahead of current workstation systems.




The speedup ratio plotted for the terapixel computation from 64 to 1024 NC6v3 V100 GPU nodes, the speedup shows close to linear scaling achieving 98% efficiency at 1024 nodes

Source: Nicolas S. Holliman, Manu Antony, James Charlton, Stephen Dowsland, Philip James, Mark Turner (Feb 13, 2019) - <https://arxiv.org/abs/1902.04820>

“Petascade Cloud Supercomputing for Terapixel Visualization of a Digital Twin”

Azure Delivers Best-in-Class Performance




CONCLUSION

“We benchmarked the performance of the best available computing hardware from public cloud providers with high performance Linpack. We optimized the benchmark for each computing environment and evaluated the relative performance for distributed memory calculations. **We found Microsoft Azure to deliver the best results, and demonstrated that the performance per single computing core on public cloud to be comparable to modern traditional supercomputing systems.**”

Source: Mohammadi, M., Bazhirov, T. (Feb 9, 2017) Exabyte Inc. “Comparative benchmarking of cloud computing vendors with High Performance Linpack” (<https://arxiv.org/pdf/1702.02968.pdf>)

How we finish meetings about HPC after our presentations?




The key elements

Key points on HPC



STORAGE



NETWORKING



COMPUTING



MANAGEMENT

Storage



Distributed network file systems

“In computing, a distributed file system (DFS) or network file system is any file system that allows access to files from multiple hosts sharing via a computer network. This makes it possible for multiple users on multiple machines to share files and storage resources”

https://en.wikipedia.org/wiki/Comparison_of_distributed_file_systems




IBM
Spectrum
Scale



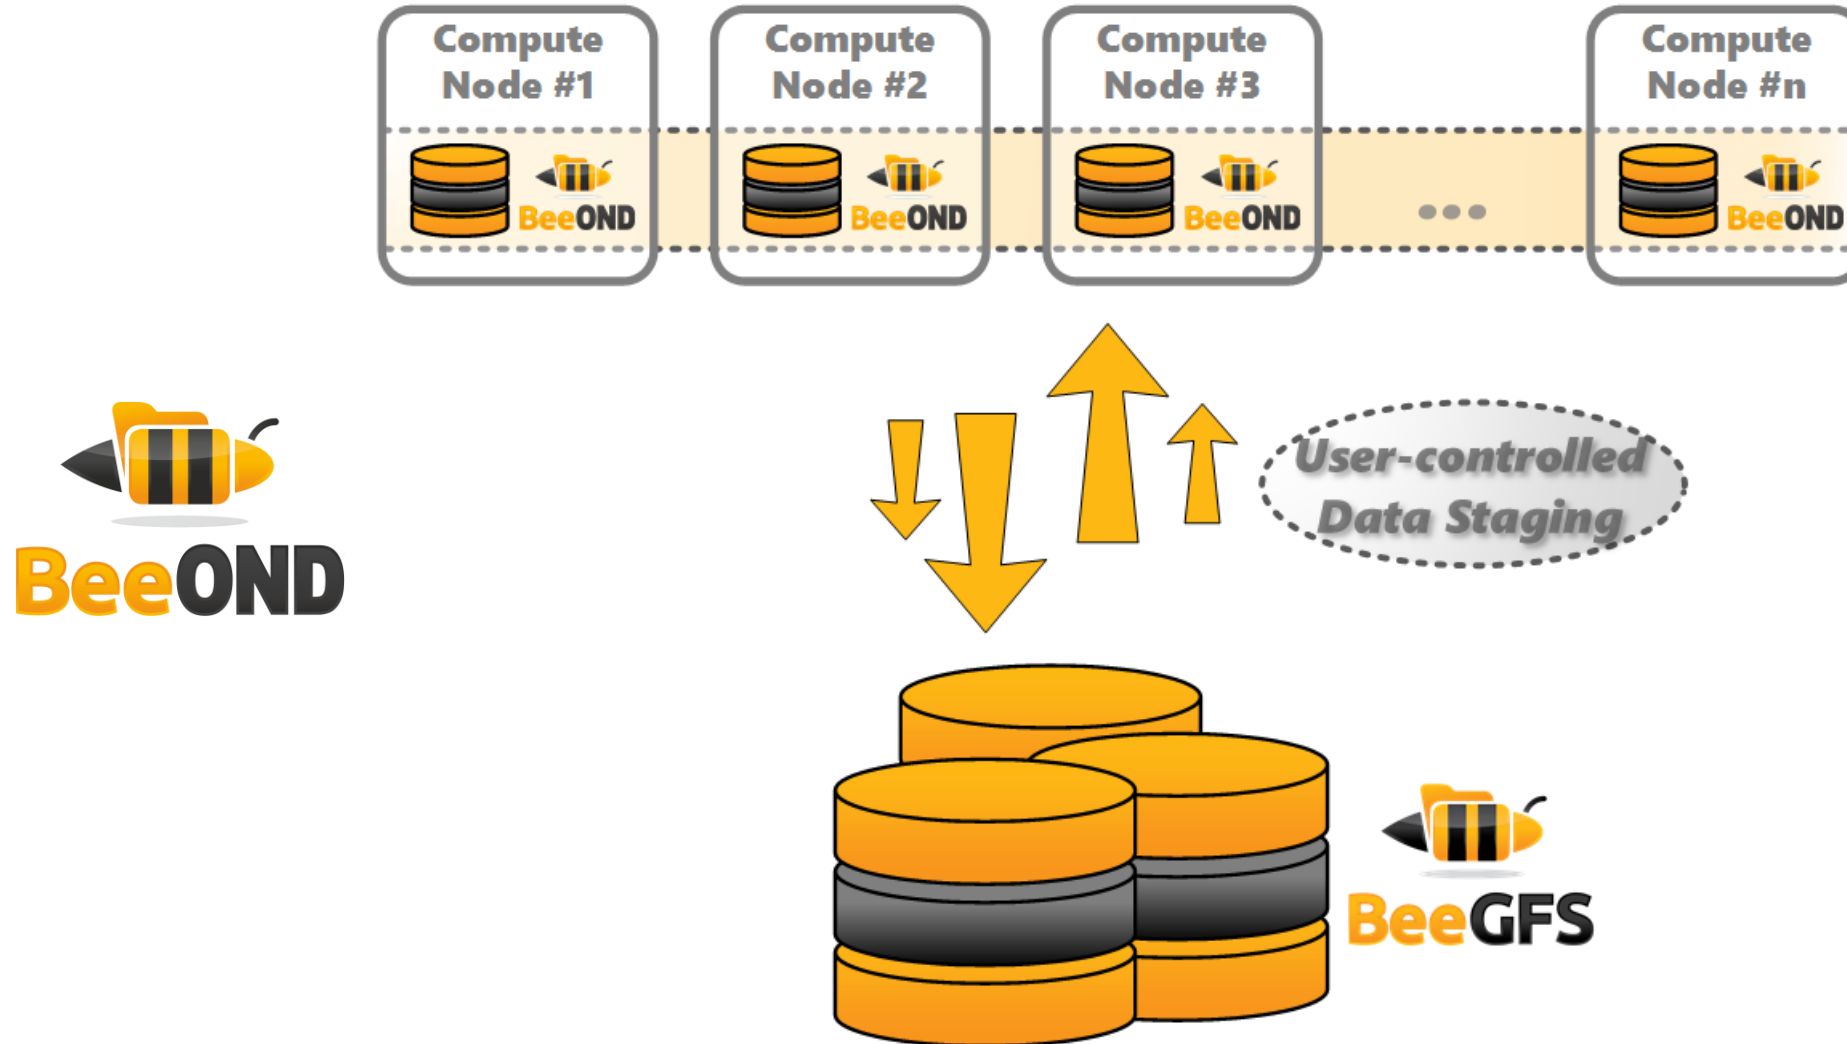
Azure NetApp Files



Distributed network file systems: an example



On-demand file system per job



Cloud-based cache layer, Avere

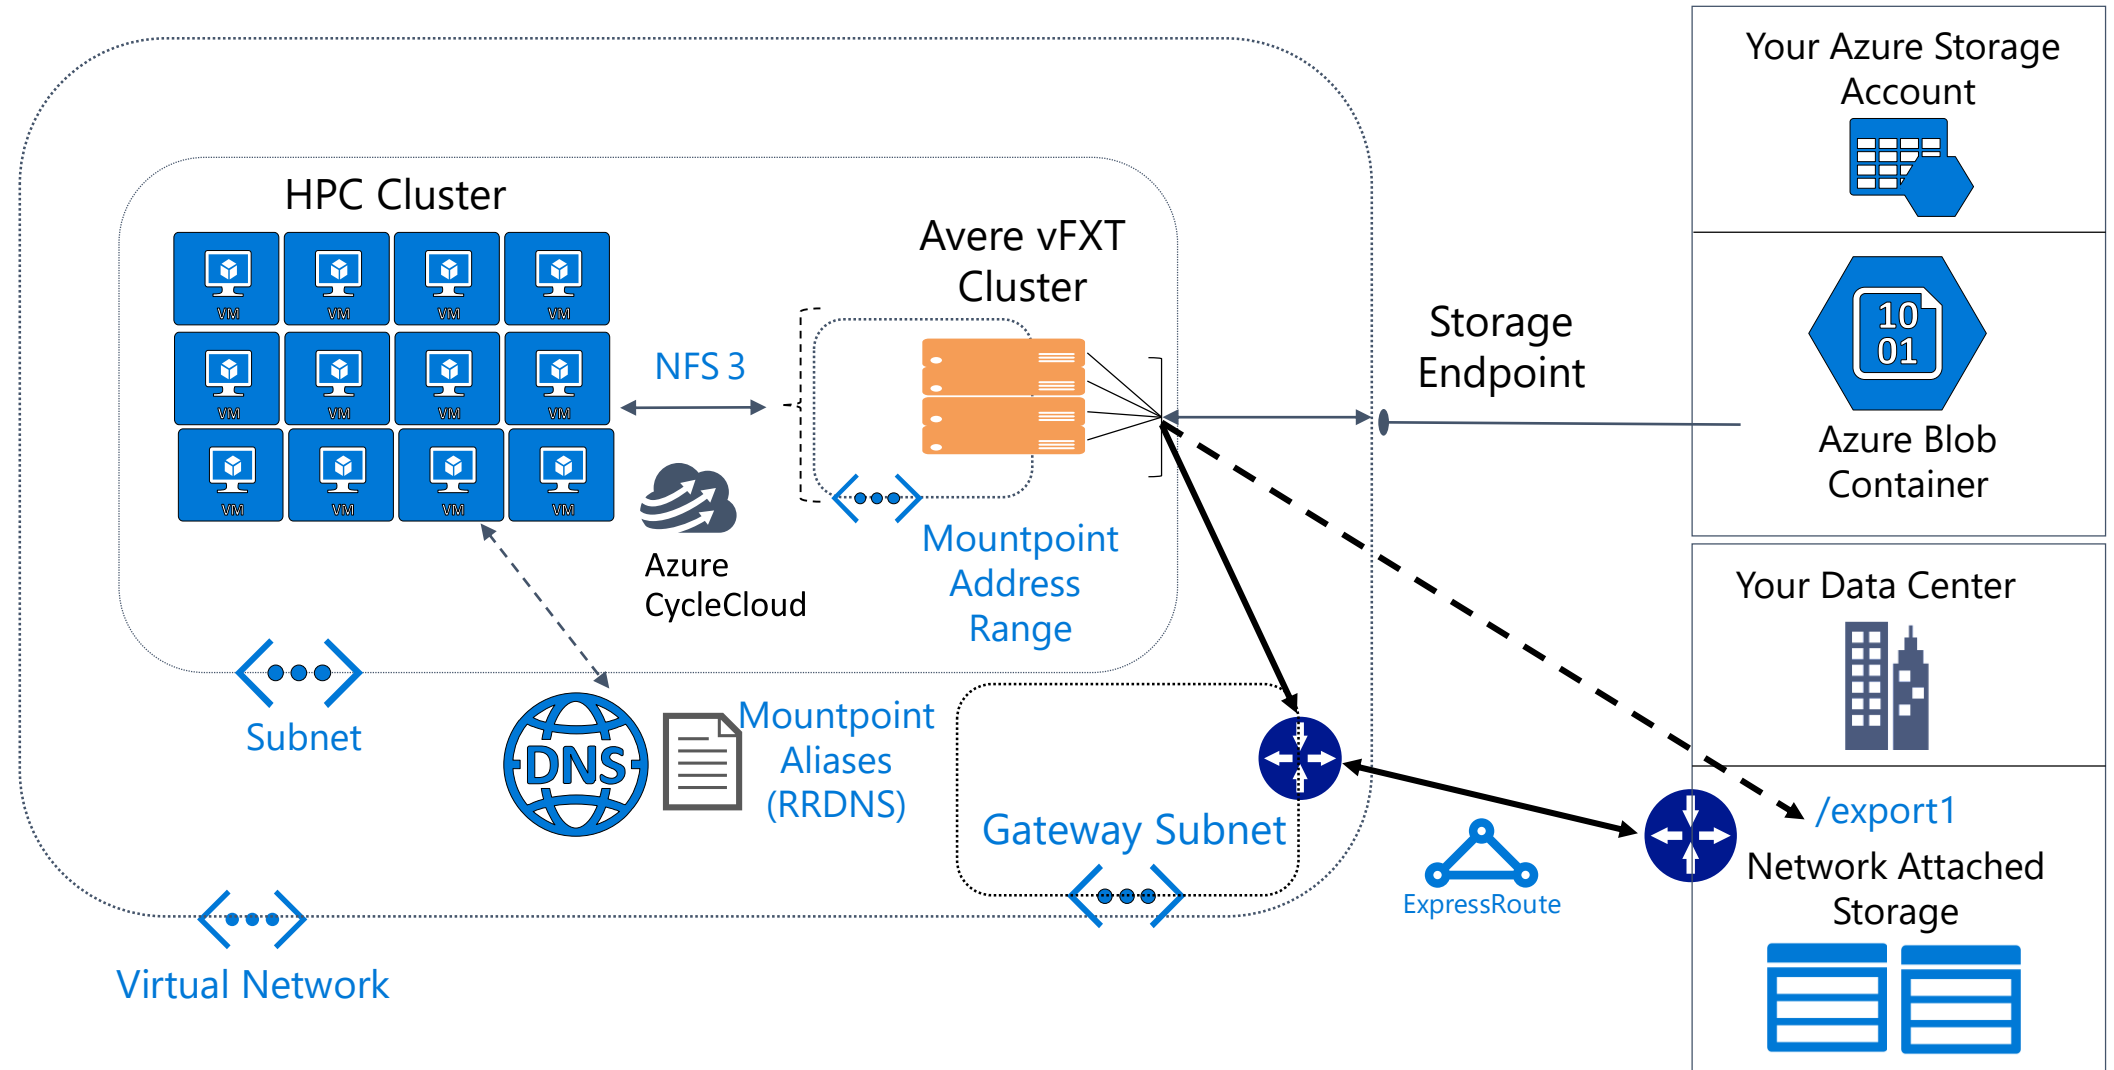


*Avere vFXT for Azure provides a **cloud-based cache layer** for data-intensive high performance computing (HPC) tasks. By caching files close to your compute nodes, Avere vFXT speeds read times and lets you work more smoothly even at peak load. Avere vFXT works best with systems that have **between 1,000 and 40,000 client cores**.*

Avere vFXT can work with your existing on-premises data storage to provide Azure-based computing resources local access to active files that are stored long-term in your datacenter.

<https://docs.microsoft.com/en-us/azure/avere-vfxt/>

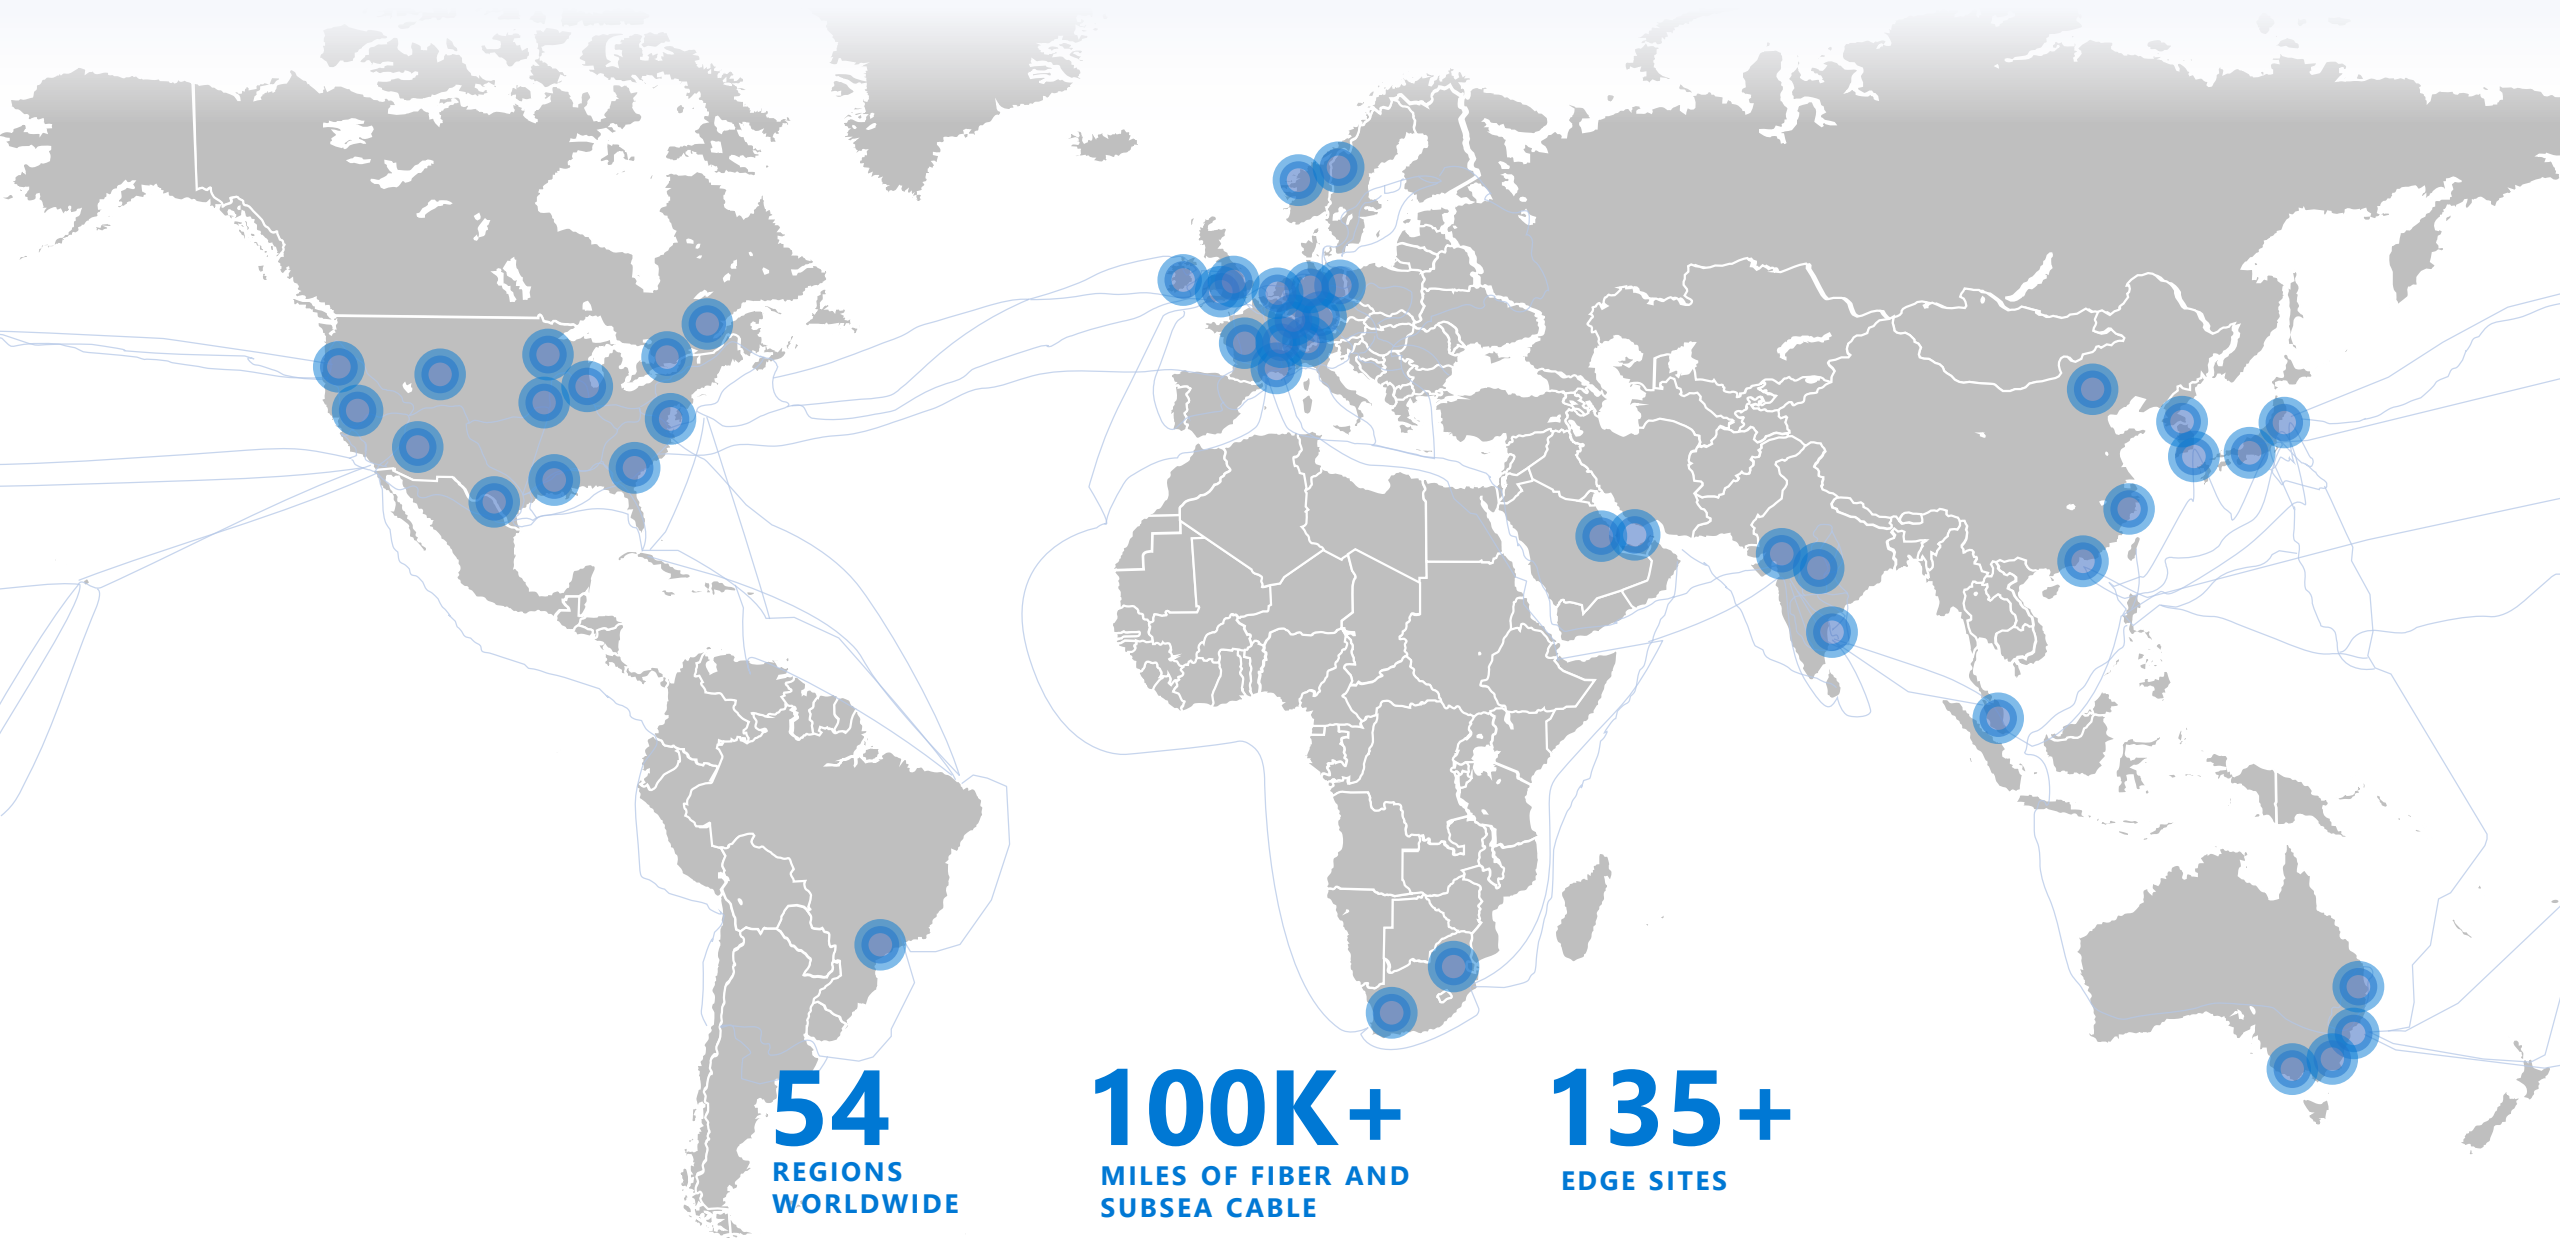
Avere architecture deployment



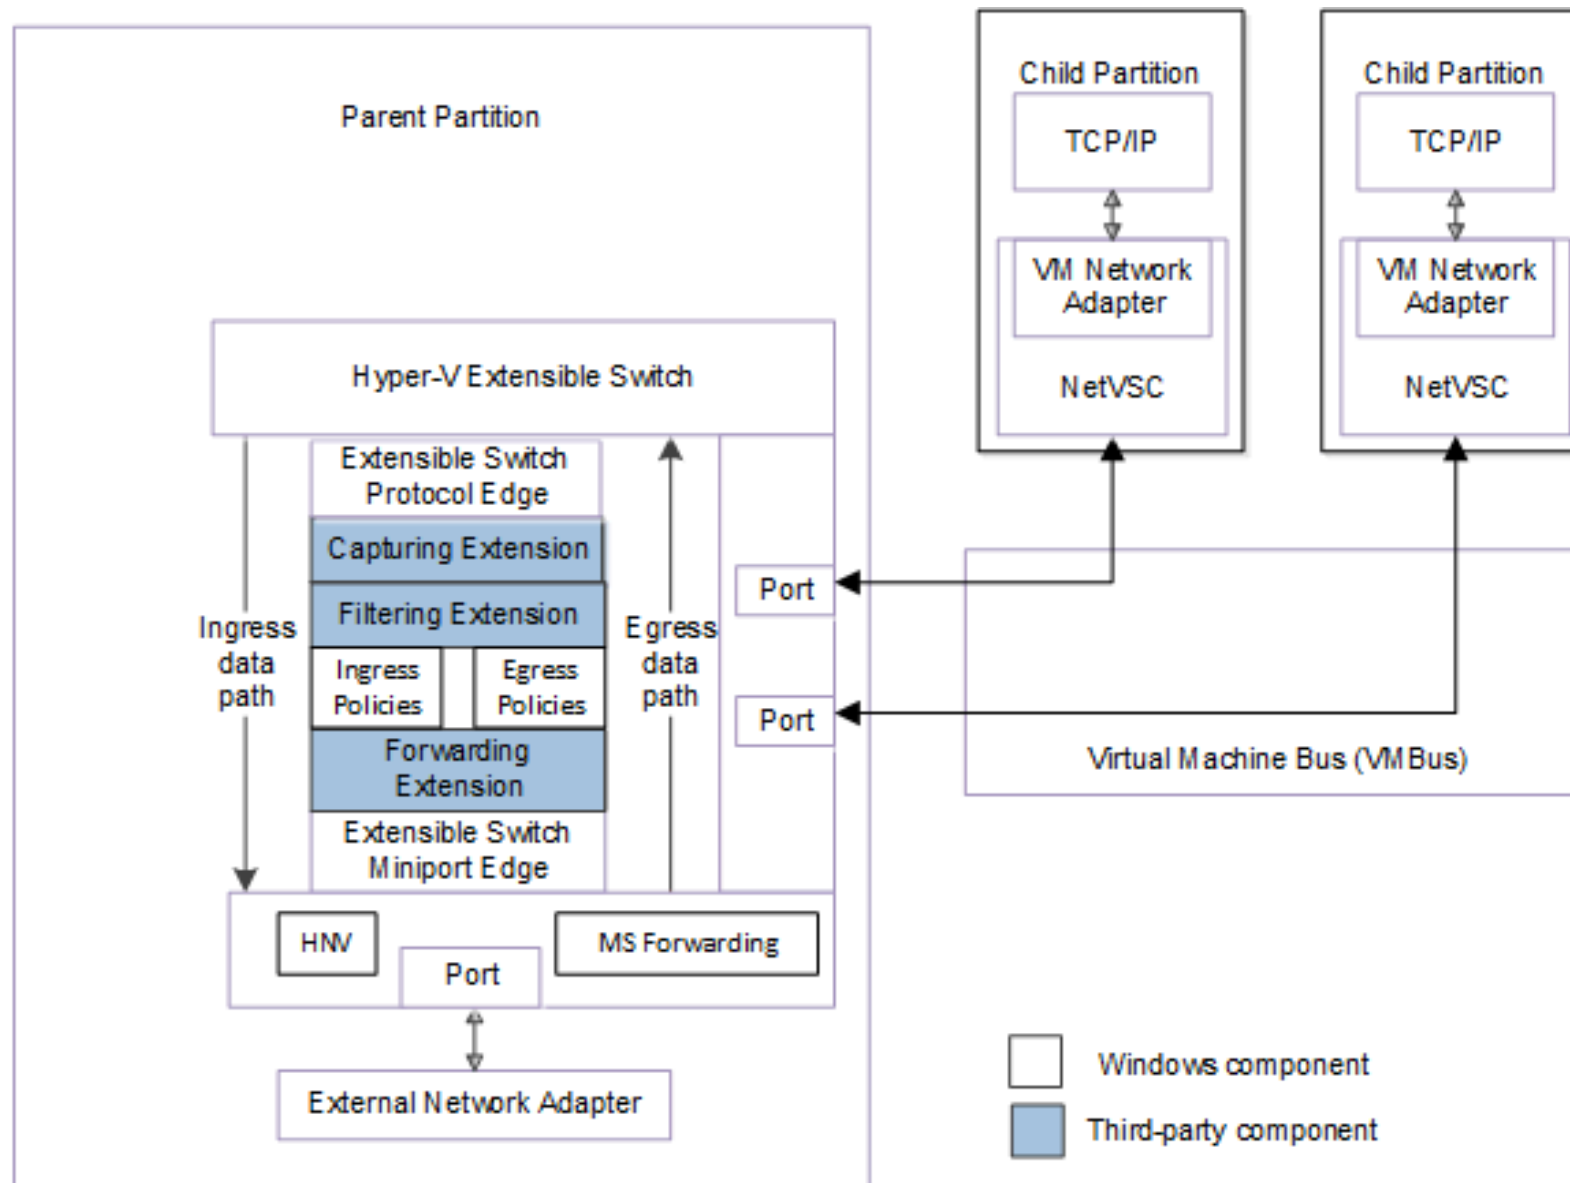
Networking

The background of the slide is a solid dark red color. Overlaid on this background is a complex, abstract network diagram. This diagram consists of numerous small, light-red circular nodes of varying sizes. These nodes are interconnected by a dense web of thin, light-red lines. Some of these lines are straight, while others are curved or wavy. The overall pattern of the network is organic and sprawling, with several distinct clusters or hubs where a large number of lines converge. The lines and nodes are most concentrated in the upper right and lower right areas, with more sparse connections towards the left side of the image.

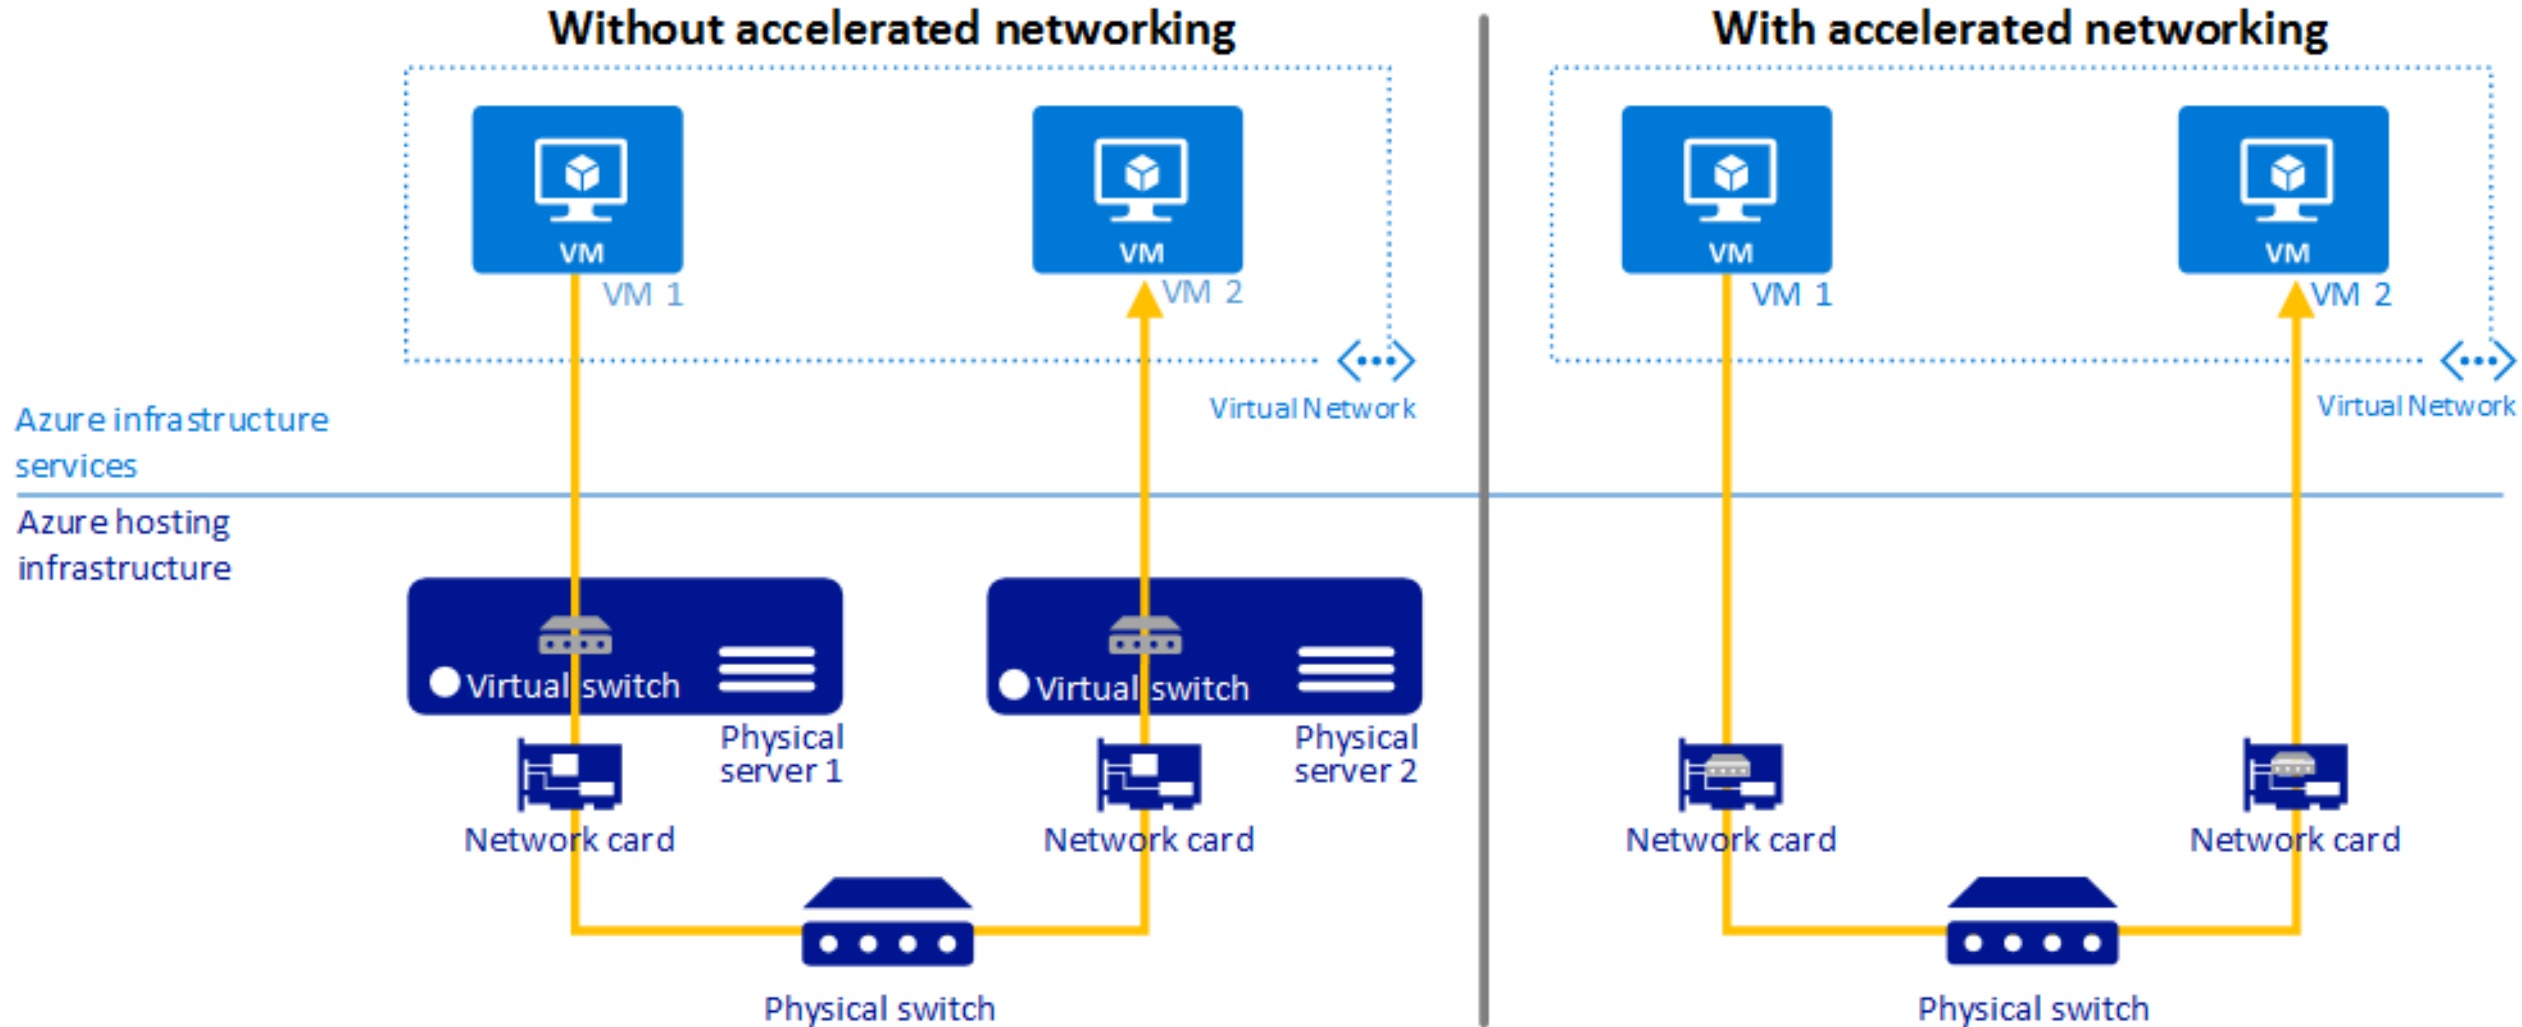
Azure Networking Infrastructure



How Hyper-v net adapter works for “dummies”?



Accelerated Networking



Infiniband....


Ethernet

vs

Infiniband



Why Infiniband matters?

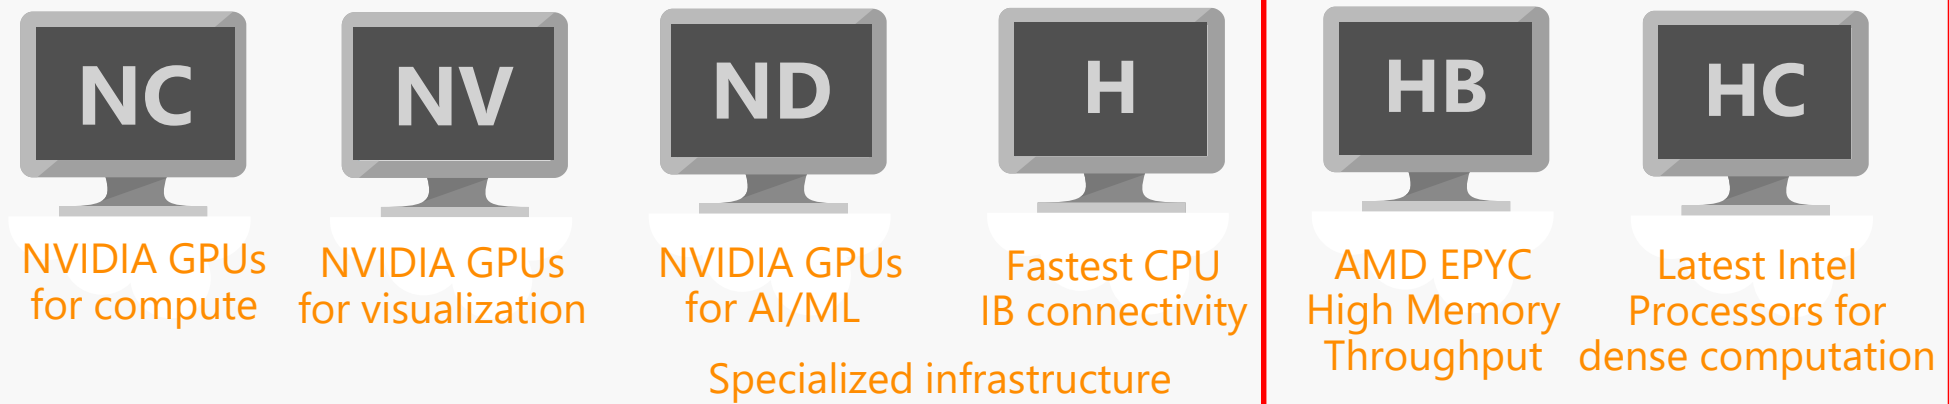


- ① Lower run time is better
- ② As you add more cores the job should run faster
- ③ Because of lack of RDMA InfiniBand low latency networking jobs run slower after 64 cores on other clouds
- ① Azure continues to scale well after adding more cores

Computing

The background is a solid red color with a complex, abstract pattern of white and light red lines, dots, and geometric shapes. The pattern includes a grid of squares in the upper left, a series of concentric circles and dots in the lower left, and a large, intricate network of lines and dots on the right side, resembling a neural network or a complex data structure. The overall effect is a high-tech, digital aesthetic.



Computing options in Azure



Management

The background is a solid red color with a complex, abstract pattern of white and light red lines, dots, and geometric shapes. The pattern includes a grid of squares in the upper left, a series of concentric circles and dots in the lower left, and a large, intricate network of lines and dots in the upper right. The overall effect is a sense of dynamic, interconnectedness.

Management



Let's Play

iGracias

Patrocinadores Locales



Colabora



2019

Global **Azure**
BOOTCAMP

